

JOVAN KALYANGO

Data Engineer — Remote

Bologna, Italy (CET) · +39 349 433 3088 · kalyangojonn@gmail.com · jovan-portfolio-chi.vercel.app · linkedin.com/in/jovan-kalyango · github.com/KalyangoJovan

PROFILE

Data Engineer with 4+ years of production experience on Google Cloud, holding an MSc in Statistical Sciences (Data Science) from the University of Bologna and Utrecht University. Daily stack: Python, SQL, BigQuery, Dataflow (Apache Beam), and Pub/Sub, building streaming and batch pipelines for enterprise analytics. Alongside consultancy work, I design, ship, and operate my own production systems end to end — a government-scale assessment platform on PostgreSQL and a fully automated market-data pipeline processing a full year of trading activity — so I am comfortable owning a data product from schema design to deployment and daily operation, remotely and autonomously.

PROFESSIONAL EXPERIENCE

Data Engineer — Bitbang Srl

Bologna, Italy · June 2022 – Present

Digital analytics consultancy serving enterprise clients

- Build and operate streaming and batch data pipelines on GCP: Pub/Sub ingestion, Apache Beam (Python SDK) transforms on Dataflow — ParDo/DoFn, windowing, dead-letter side outputs — landing in BigQuery via the Storage Write API and write-disposition patterns.
- Develop analysis-ready BigQuery models with advanced SQL: window functions (deduplication via ROW_NUMBER/QUALIFY), MERGE upserts, partitioning and clustering for cost control, and scheduled queries for recurring delivery.
- Translate enterprise clients' business questions into reliable, automated ETL/ELT workflows, reducing manual processing and improving data reliability for reporting and BI.

Founder & Engineer — Rubikon (therubikon.org)

Remote · 2025 – Present

Governance and competency assessment SaaS — live in production

- Designed and shipped a production multi-tenant assessment platform (Next.js 15, Prisma, Neon PostgreSQL, Vercel) used in a government ministry engagement covering 652 officers across 4 grade cohorts.
- Built a hybrid paper + digital data-ingestion pipeline: CSV template generation, two-step preview→commit upload with row-level validation, unified storage so paper and digital responses analyse together; plus gap-analysis heatmaps and wide-format CSV export.
- Applied additive-only, transaction-wrapped schema changes against a live production database with before/after row-count verification — zero downtime, zero data loss.
- Automated branded client deliverables: server-side Word report generation (docx) producing individual and organisational assessment reports from raw response data.

Founder — Driips Capital

Remote · 2025 – Present

Systematic trading — automated research and reporting infrastructure

- Built an end-to-end automated market-data pipeline: 1-minute price-data ingestion, a Python backtesting engine validated out-of-sample, and an auto-journaling service that logged 222 trades over a full year into a reporting web cockpit with per-trade charts.
- Ran the full research lifecycle on the same infrastructure: hypothesis specification, backtest, out-of-sample confirmation, and live signal generation bridged to broker execution.

EDUCATION

MSc, Statistical Sciences (Data Science)

University of Bologna & Utrecht University

BSc, Statistical Sciences (Maths & Stats)

University of Bologna

Includes study periods at UC Berkeley and the University of Glasgow.

TECHNICAL SKILLS

Languages: Python (pandas, Apache Beam SDK), SQL (BigQuery, PostgreSQL), TypeScript

Cloud & Data: GCP: BigQuery, Dataflow, Pub/Sub, Cloud Run · Neon PostgreSQL · Prisma · streaming & batch pipeline design

Engineering: Git/GitHub, CI-safe production migrations, API design (FastAPI/Next.js routes), automated reporting (docx), Vercel deployment

Languages (spoken): English (fluent) · Italian (professional working proficiency)